

CSDS 440: Machine Learning

Soumya Ray (he/him, sray@case.edu)

Olin 516


Office hours T, Th 11:15-11:45 or by appointment

Announcements/Status

- Groups
- Cwru-courses/ Repositories

Play this 2-player game!

- Draw this board on a sheet of paper
- Each player chooses/makes three “pieces”
- Take turns placing your pieces on the intersections
- Once all pieces are placed, take turns to move them to any adjacent open intersection
- Goal: starting with empty board, first player to make three pieces in a line wins
- If no winner after 10 moves, draw
- Each person in a group should play at least 2 games (AB/AC/BC)



Today

- Foundations of Machine Learning

What is “machine learning”?


- “Machine”=autonomous system
 - No (or limited) human intervention
 - Robots, software agents, etc.

What is “Learning”?

“Learning denotes changes in the system that enable the system to do the same task more effectively the next time.”

—Herbert Simon (1916-2001)

- Also, how to do *related* tasks more effectively



A Specification for a Learning System

- Given:
 - Learning task,
 - Task examples E ,
 - Performance measure P
- Do: Produce a **concept** that is good with respect to P on *all* examples of the task
 - Measured by proxy on E

Example

- Learning Task: Learn to play a board game
- Performance measure?
 - Games won/lost
- Examples?
 - Games played (sequences of moves till win/loss)
- Concept?
 - Some function mapping current state of game to suitable moves to play

Idea

- If the learning system plays/sees enough games,
- And it produces a mapping from game state to moves (concept),
- And this concept does well with respect to the measure of “number of games won”,
- Then the system has “**learned to play the game**”

Other Examples

- Learn to recognize lions
 - E : animals, annotated “lion” or “not-lion”
 - P : fraction of animals correctly recognized as lion/not-lion
- Learn to drive
 - E : sequence of road/traffic conditions and correct vehicle operation
 - P : distance traveled without accident

Two Phases of Learning

- “Learning” or “Training” phase
 - Reason about the examples E
 - Formulate a concept that does well w.r.t. P on E
 - Could also use any prior knowledge
- “Evaluation” or “Testing” phase
 - Use learned concept on future, novel examples

Online and Batch (Offline) Learning

- **Batch/Offline** Learning: one learning phase, with a large set of examples, followed by a testing phase
- **Online** learning: Examples arrive one at a time (or in small groups); learning and evaluation phases are iterated

Inductive Generalization

- In all learning problems, need to reason from specific examples to a general case
 - Memorization \neq Learning
- Other kinds of reasoning
 - deduction (general to specific)
 - abduction (most likely cause)

Target Concept

- The unknown underlying concept that solves the learning task
 - E.g., “has-fur” and “long-teeth” and “looks-scary” → “lion”
- Typically, P will be a measure of difference between the learner’s concept and the target concept, with respect to E

Hypothesis Space

- Defines the space of general concepts the learning system will consider
 - E.g., all possible conjunctions of animal properties
 - “has-fur” and “long-teeth” and “looks-scary”, “has-fur” and “long-teeth” and NOT-“looks-scary” , “has-fur” and NOT-“long-teeth” and “looks-scary”
- Ideally, target concept is a member of this space
 - What if it isn't?
 - Maybe we should include *all possible* hypotheses?

No “Tabula Rasa” Learning

- A space that includes all possible hypotheses also
 - Contains many overly complex concepts
 - Contains the concept that memorizes E
 - Indistinguishable from target by any P (w.r.t. E)
 - *May be* too big to search feasibly
- For effective inductive generalization
 - **Must** restrict hypothesis space
 - while still (hopefully) keeping the target concept in it

Inductive Bias

- The set of assumptions used by a learning system to restrict its hypothesis space
- The more assumptions made, the “stronger” the bias
- Can quantify this (later)