

CSDS 440: Machine Learning

Soumya Ray (he/him, sray@case.edu)

Olin 516

Office hours T, Th 11:15-11:45 or by appointment

Announcements

- Quiz next Tuesday
 - Topics up to and including today's lecture
 - Bring calculator

Ensemble Methods

- So far, for each task, we construct a single classifier
- Suppose we had a way of constructing *multiple* classifiers and *combining their output*
 - Would this generalize any better than constructing a single classifier?
 - Why or why not?
 - And how do we construct multiple classifiers in the first place?

Single vs. multiple classifiers

- Suppose for some problem we have k classifiers h_1, \dots, h_k that:
 - Each has error less than chance: $\varepsilon_i < 1/2$
 - Make *uncorrelated* errors on new examples
- Suppose we combine their predictions on a new example via majority vote
- What is the error rate of the combined system?


Single vs. multiple classifiers

- Since we combine using majority voting, to get an example wrong, at least half the classifiers must be wrong
- Since the errors of each classifier is uncorrelated, we can treat each classifier as an “independent trial” for the new example

Single vs. multiple classifiers

- Assume all the classifiers have the same error rate ε
 - No loss of generality—can choose max error over all classifiers to get upper bound on error rate of combined system
- Then the number of wrong classifiers follows a Binomial distribution with parameters ε, k

Example: $\varepsilon=0.3, k=21$



$\Pr(\text{Combined System is wrong}) =$
 $\Pr(\text{At least 11 classifiers are wrong}) =$
0.026

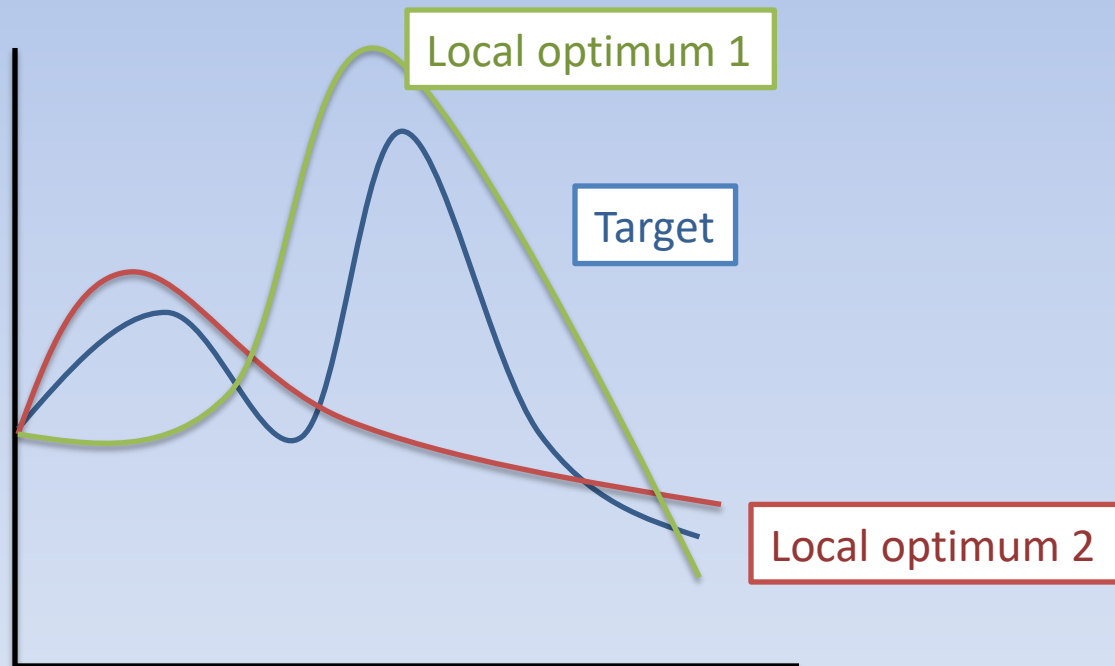
“Ensembles” of Classifiers

- A collection of classifiers combined using some sort of voting scheme is called an “ensemble”
- Getting classifiers that have error less than chance is (usually) easy
 - Note that this is *generalization* error
- Getting classifiers that make uncorrelated errors is usually not
 - But even if not, ensembles can outperform single classifiers in practice

Why ensembles do well in practice (1)

- Many classifiers are a result of some search procedure (e.g. gradient descent) which can get stuck in local optima
- Averaging these “local optimum” classifiers can provide a better approximation to the target function

Picture



Why ensembles do well in practice (2)

- An ensemble of classifiers may have a more complex decision boundary than any single classifier
 - E.g. an ensemble of voting linear classifiers is not (generally) a linear classifier

Why ensembles do well in practice (3)

- Consider probabilistic classification
- Given a training sample D , what is the most probable classification for a new instance \mathbf{x}_{new} ?
 - i.e. what is $\Pr(Y_{new} = y \mid D, \mathbf{x}_{new})$?
 - Assume you are investigating a hypothesis class H

Bayesian Model Averaging

$$\begin{aligned}\Pr(Y_{new} = y \mid D, \mathbf{x}_{new}) \\&= \sum_{h \in H} \Pr(Y_{new} = y \mid D, h, \mathbf{x}_{new}) \Pr(h \mid D, \mathbf{x}_{new}) \\&= \sum_{h \in H} \Pr(Y_{new} = y \mid h, \mathbf{x}_{new}) \Pr(h \mid D)\end{aligned}$$

Classification according
to h

Posterior probability
of h given training sample

- Also called Bayes optimal classification
- Cannot be outperformed on average by any single hypothesis in H

The downsides

- How large an ensemble to use is not well understood
 - Too small—no effect, too large—tends to overfit
- An ensemble is usually much harder to interpret than a single classifier
 - E.g. a decision tree vs a set of trees
- Computation time, memory etc. all increase
 - (sometimes we can parallelize)

General Ensemble Construction

- Can construct ensembles in several ways
 - Modifying the training set
 - Modifying the set of attributes
 - Modifying the outputs
 - Randomizing the learning algorithm


Modifying the Training Set

- General idea:
 - Create multiple training sets, each different from the others in some way
 - Apply learning algorithm to each set
 - Resulting classifiers vote on new examples
- Works best for “unstable” algorithms
 - Small change to data can lead to large change in solution
- Two important methods
 - Bagging
 - Boosting

Bagging (BREIMAN 96)

- “**Bootstrap Aggregation**”
- Each training sample is a bootstrap replicate of the initial set
 - If the set has size m , sample m examples *uniformly with replacement* from it
- To classify a new example, use majority voting

Picture



Bagging

- In practice, Bagging:
 - Rarely or never hurts accuracy
 - But improvements in accuracy are likewise small
- Voting classifiers constructed with bootstrap replicates can result in “averaging out” the effect of noise