

CSDS 440: Machine Learning

Soumya Ray (he/him, sray@case.edu)

Olin 516

Office hours T, Th 11:15-11:45 or by appointment

Support Vector Machines

- Combines three fundamental ideas
 - Linear discriminants
 - Margins
 - Kernels
- A theoretically well justified and empirically well-behaved method arising from three fields: ML (Cortes & Vapnik), Statistics (Wahba), Operations Research (Mangasarian)

The primal (separable) SVM

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{so that } \forall i, -[y_i(\mathbf{w} \bullet \mathbf{x}_i + b) - 1] \leq 0$$

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{so that } \forall i, -[y_i(\mathbf{w} \bullet \varphi(\mathbf{x}_i) + b) - 1] \leq 0$$

The primal (separable) SVM

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{so that } \forall i, -[y_i(\mathbf{w} \bullet \mathbf{x}_i + b) - 1] \leq 0$$

$$\therefore \ell(\mathbf{w}, b, \boldsymbol{\alpha}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_i \alpha_i [y_i(\mathbf{w} \bullet \mathbf{x}_i + b) - 1]$$

Linearly-separable SVM, Dual Form

$$\max_{\alpha} D(\alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j \mathbf{x}_i \cdot \mathbf{x}_j$$

so that $\alpha \geq 0$, $\sum_i \alpha_i y_i = 0$



From derivative w.r.t b

Kernels

- Notice that by using the dual formulation, we could write the formulation and the solution in terms of $\mathbf{x}_i \cdot \mathbf{x}_j$
 - In general, $\varphi(\mathbf{x}_i) \cdot \varphi(\mathbf{x}_j)$
- Define the **kernel** to be $K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i) \cdot \varphi(\mathbf{x}_j)$
- For m examples, get an $m \times m$ matrix --- the **kernel matrix**

Nonlinear SVM, kernelized dual form

$$\max_{\alpha} D(\alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j K(\mathbf{x}_i, \mathbf{x}_j)$$

$$\text{so that } \alpha \geq 0, \sum_i \alpha_i y_i = 0$$

Why is this useful?

- Given ϕ , easy to find K
- More interestingly, for certain functions K , can show *there must exist a ϕ* for which K is a kernel
- This ϕ could be very high dimensional, but the kernel computation is much more efficient
- This allows us to do classification in very high dimensional spaces, without ever “mapping” the data into those spaces



“Kernel trick”

Example

$$K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} \cdot \mathbf{b})^2 \leftarrow O(n) \text{ computation}$$

$$= \left(\sum_i a_i b_i \right) \left(\sum_j a_j b_j \right)$$

$$= \sum_{i,j} (a_i a_j)(b_i b_j)$$

$$= \varphi(\mathbf{a}) \cdot \varphi(\mathbf{b}), \text{ where } \leftarrow O(n^2) \text{ computation!}$$

$$\varphi(\mathbf{x}) = [x_1^2, \sqrt{2}x_1x_2, \dots, \sqrt{2}x_ix_j, \dots, x_n^2]$$

Example

$$\mathbf{a} = (1, 2), \mathbf{b} = (3, 4)$$

$$K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} \cdot \mathbf{b})^2 = (1 \times 3 + 2 \times 4)^2 = 121$$

If $K(\mathbf{a}, \mathbf{b}) = \varphi(\mathbf{a})\varphi(\mathbf{b}) = (\mathbf{a} \cdot \mathbf{b})^2$ then

$$\varphi(\mathbf{a}) = [a_1^2, \sqrt{2}a_1a_2, a_2^2] = [1, 2\sqrt{2}, 4]$$

$$\varphi(\mathbf{b}) = [b_1^2, \sqrt{2}b_1b_2, b_2^2] = [9, 12\sqrt{2}, 16]$$

$$\varphi(\mathbf{a})\varphi(\mathbf{b}) = 9 + 48 + 64 = 121$$

What is a valid kernel?

- Intuitively, the dot product measures the (unnormalized) cosine of the angle between two vectors
 - A measure of similarity
 - “large” $K(\mathbf{x}, \mathbf{y}) \rightarrow \mathbf{x}$ and \mathbf{y} are “similar”
- Suppose we choose some other function that measures similarity
 - E.g., $K(\mathbf{x}, \mathbf{y}) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{y}\|^2}{2\sigma^2}\right)$
 - Is this a valid kernel?

Yes! Called a **Gaussian or RBF Kernel**.
Corresponds to infinite-dimensional ϕ !!

Mercer's Conditions

Let $K : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ be a function. K is a valid kernel iff for all finite $\{\mathbf{x}_1, \dots, \mathbf{x}_m\}$, the kernel matrix is symmetric positive semidefinite.

Symmetry: $K(\mathbf{x}, \mathbf{y}) = K(\mathbf{y}, \mathbf{x})$

PSD ($K \geq 0$): $\forall \mathbf{v} \neq 0, \mathbf{v}^T K \mathbf{v} \geq 0$



Necessary *and* sufficient!

- Given any kernel(s), can compose them in various ways to get other kernels

Classification

$$\max_{\alpha} D(\alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j K(\mathbf{x}_i, \mathbf{x}_j)$$

$$\text{so that } \alpha \geq 0, \sum_i \alpha_i y_i = 0$$

- Note that classification also does not require φ :

$$\mathbf{w} \bullet \varphi(\mathbf{x}_{new}) = \sum_{i \in \text{Support Vectors}} \alpha_i y_i \varphi(\mathbf{x}_i) \bullet \varphi(\mathbf{x}_{new})$$

$$= \sum_{i \in \text{Support Vectors}} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}_{new})$$

Generalizing kernels

- Kernel functions can be used in many other contexts, thanks to the **Representer Theorem**

Representer Theorem

- Any optimization program of the form

$$\min_f \frac{1}{2} g(\|f\|) + C \sum_i L(y_i, f(\varphi(\mathbf{x}_i)))$$

- With g monotonic has a solution that looks like:

$$f(\mathbf{x}) = \sum_i \alpha_i y_i K(\mathbf{x}, \mathbf{x}_i)$$

- (Kimeldorf and Wahba, Scholkopf et al)

SVM

- Standard SVM:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \xi_i$$

so that $\forall i, [y_i(\mathbf{w} \bullet \mathbf{x}_i + b) + \xi_i] \geq 1, \xi_i \geq 0$


So $\xi_i = (1 - y_i(\mathbf{w} \bullet \mathbf{x}_i + b))_+$

Plus function (Also RELU)

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i [(1 - y_i(\mathbf{w} \bullet \mathbf{x}_i + b))_+]$$

(Hinge Loss)

“Hinge” Loss



Logistic Regression

- Let us look again at the objective function we optimize to get LR (with overfitting control):

$$\mathbf{w}, b = \arg \min \frac{1}{2} \|\mathbf{w}\|^2 + C \left[\sum_{i \in pos} -\log(P(y = 1 | \mathbf{x}_i)) + \sum_{i \in neg} -\log(1 - P(y = 1 | \mathbf{x}_i)) \right]$$

Rewriting the LR objective

- Note that, if we make the class labels 1 and -1:

$$p(Y = 1 \mid \mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}$$

$$p(Y = -1 \mid \mathbf{x}) = 1 - \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} = \frac{e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} = \frac{1}{1 + e^{(\mathbf{w} \cdot \mathbf{x} + b)}}$$

$$\text{So } p(Y = y \mid \mathbf{x}) = \frac{1}{1 + e^{-y(\mathbf{w} \cdot \mathbf{x} + b)}}, \text{ and}$$

$$-\log p(Y = y \mid \mathbf{x}) = \log \left(1 + e^{-y(\mathbf{w} \cdot \mathbf{x} + b)} \right)$$

Rewriting the LR objective

- So we can rewrite the objective:

$$\mathbf{w}, b = \arg \min \frac{1}{2} \|\mathbf{w}\|^2 + C \left[\sum_{i \in pos} -\log(P(y = 1 | \mathbf{x}_i)) + \sum_{i \in neg} -\log(1 - P(y = 1 | \mathbf{x}_i)) \right]$$
$$= \arg \min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \log(1 + e^{-y_i(\mathbf{w} \cdot \mathbf{x}_i + b)})$$

SVM and Logistic Regression

- Standard SVM:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \left[(1 - y_i (\mathbf{w} \cdot \mathbf{x}_i + b))_+ \right]$$

Hinge Loss

Margin


Error on training data

- (Regularized) Logistic Regression:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \log \left(1 + e^{-y_i (\mathbf{w} \cdot \mathbf{x}_i + b)} \right)$$

Logistic Loss

“Hinge” Loss vs Logistic Loss



Kernel Logistic Regression

- The LR objective “looks like” an SVM with an alternative loss function
- By the Representer theorem, the solution to this also must be $\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$

- And so if we introduce kernels:

$$f(\mathbf{x}) = \sum_i \alpha_i y_i K(\mathbf{x}, \mathbf{x}_i)$$

Zhu and Hastie, “KLR and the Import Vector Machine” (ask for paper)

Pros and Cons of SVMs

- + Well-justified, rigorous theory combines fundamental ideas
- + “Builds” representation (connection to deep learning); can learn very complex hypotheses; has “built-in” overfitting control
- + Numerous practical applications; often very good performance
- + Kernelizing is a very powerful idea thanks to the representer theorem
- Can be sensitive to noise and outliers with nonlinear kernels
- Requires input scaling
- Not so easy to implement
- Very memory intensive due to large kernel matrices and constraints
- Not easy to parallelize/GPUize